

# Lecture 3: Linear Classification

# Last time: Image Classification

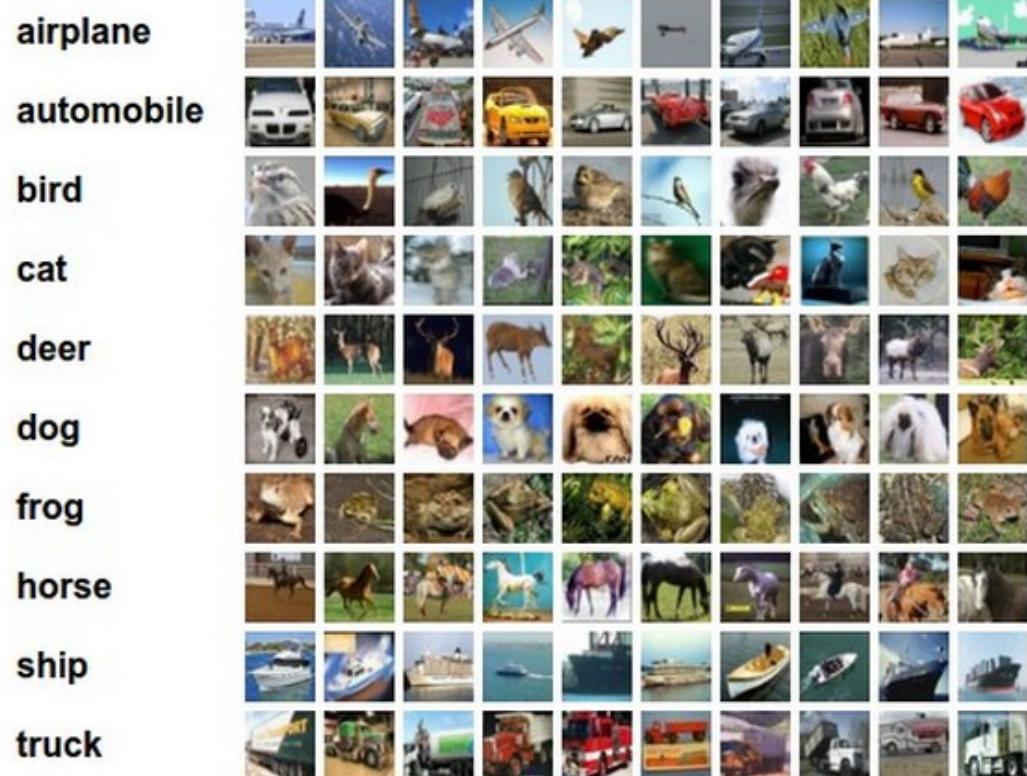


assume given set of discrete labels  
{dog, cat, truck, plane, ...}

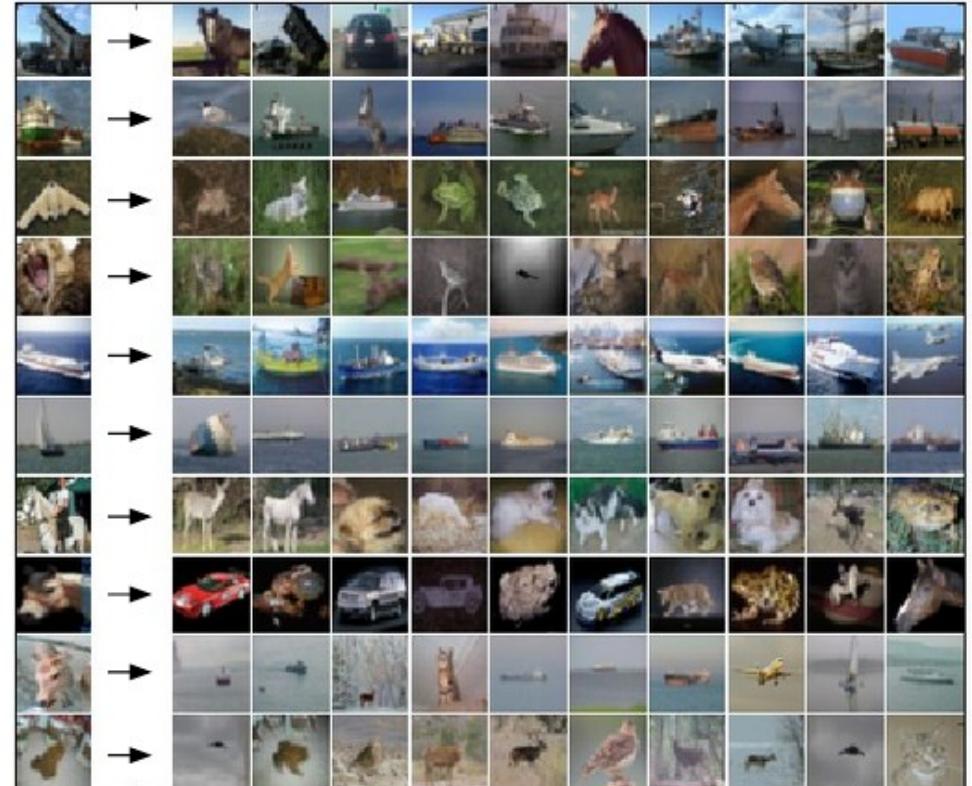
→ cat

# k-Nearest Neighbor

training set



test images



# Linear Classification

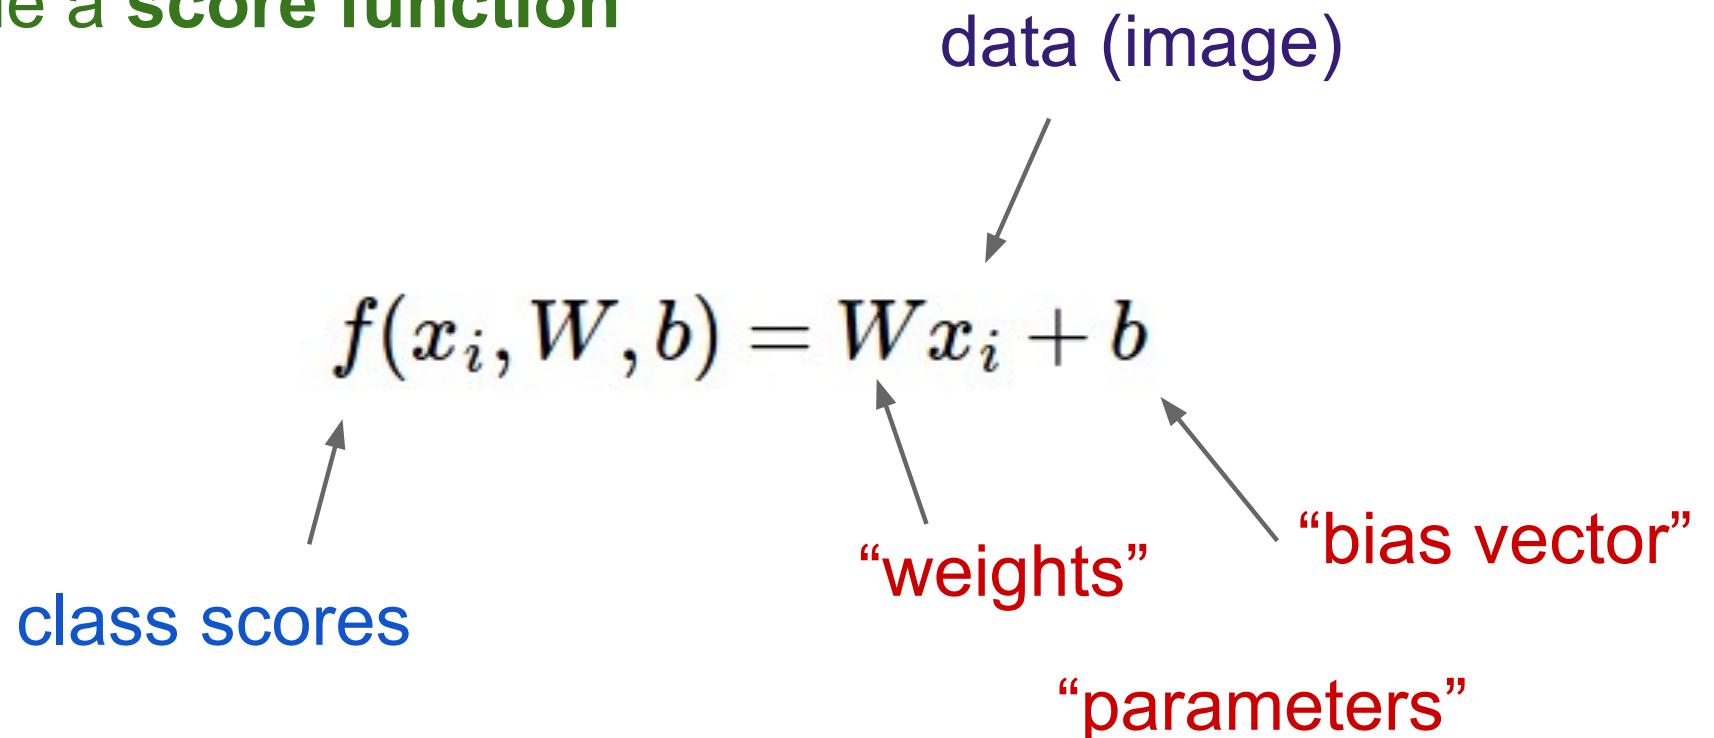
1. define a **score function**



class scores

# Linear Classification

## 1. define a **score function**



# Linear Classification

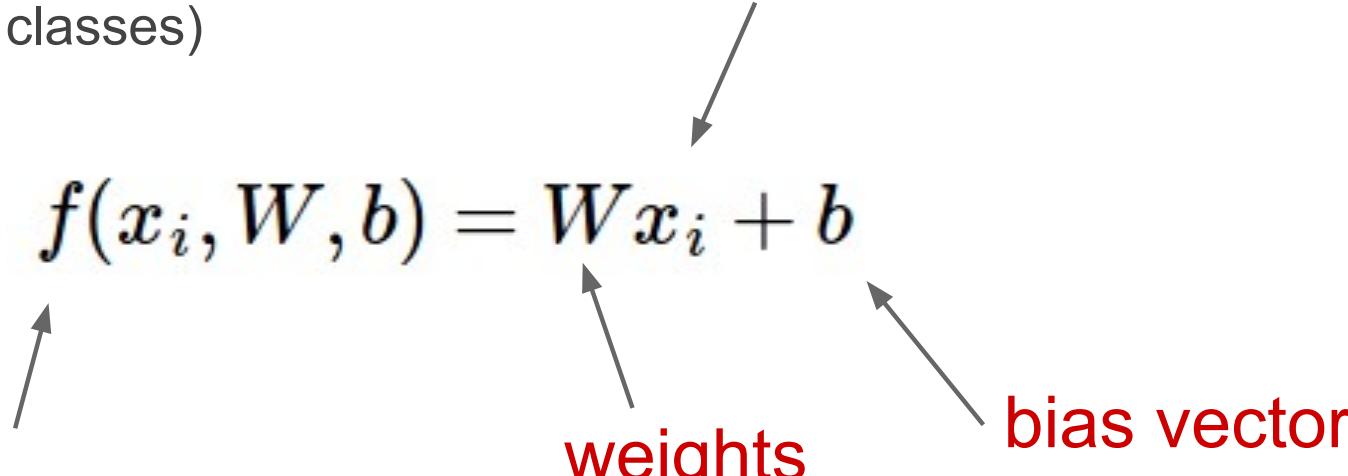
## 1. define a **score function**

(assume CIFAR-10 example so  
32 x 32 x 3 images, 10 classes)

data (image)  
[3072 x 1]

$$f(x_i, W, b) = Wx_i + b$$

class scores                      weights                      bias vector



# Linear Classification

## 1. define a **score function**

(assume CIFAR-10 example so  
32 x 32 x 3 images, 10 classes)

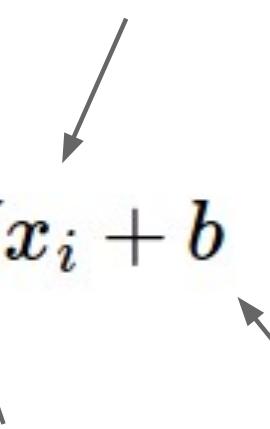
class scores  
[10 x 1]

$$f(x_i, W, b) = Wx_i + b$$

data (image)  
[3072 x 1]

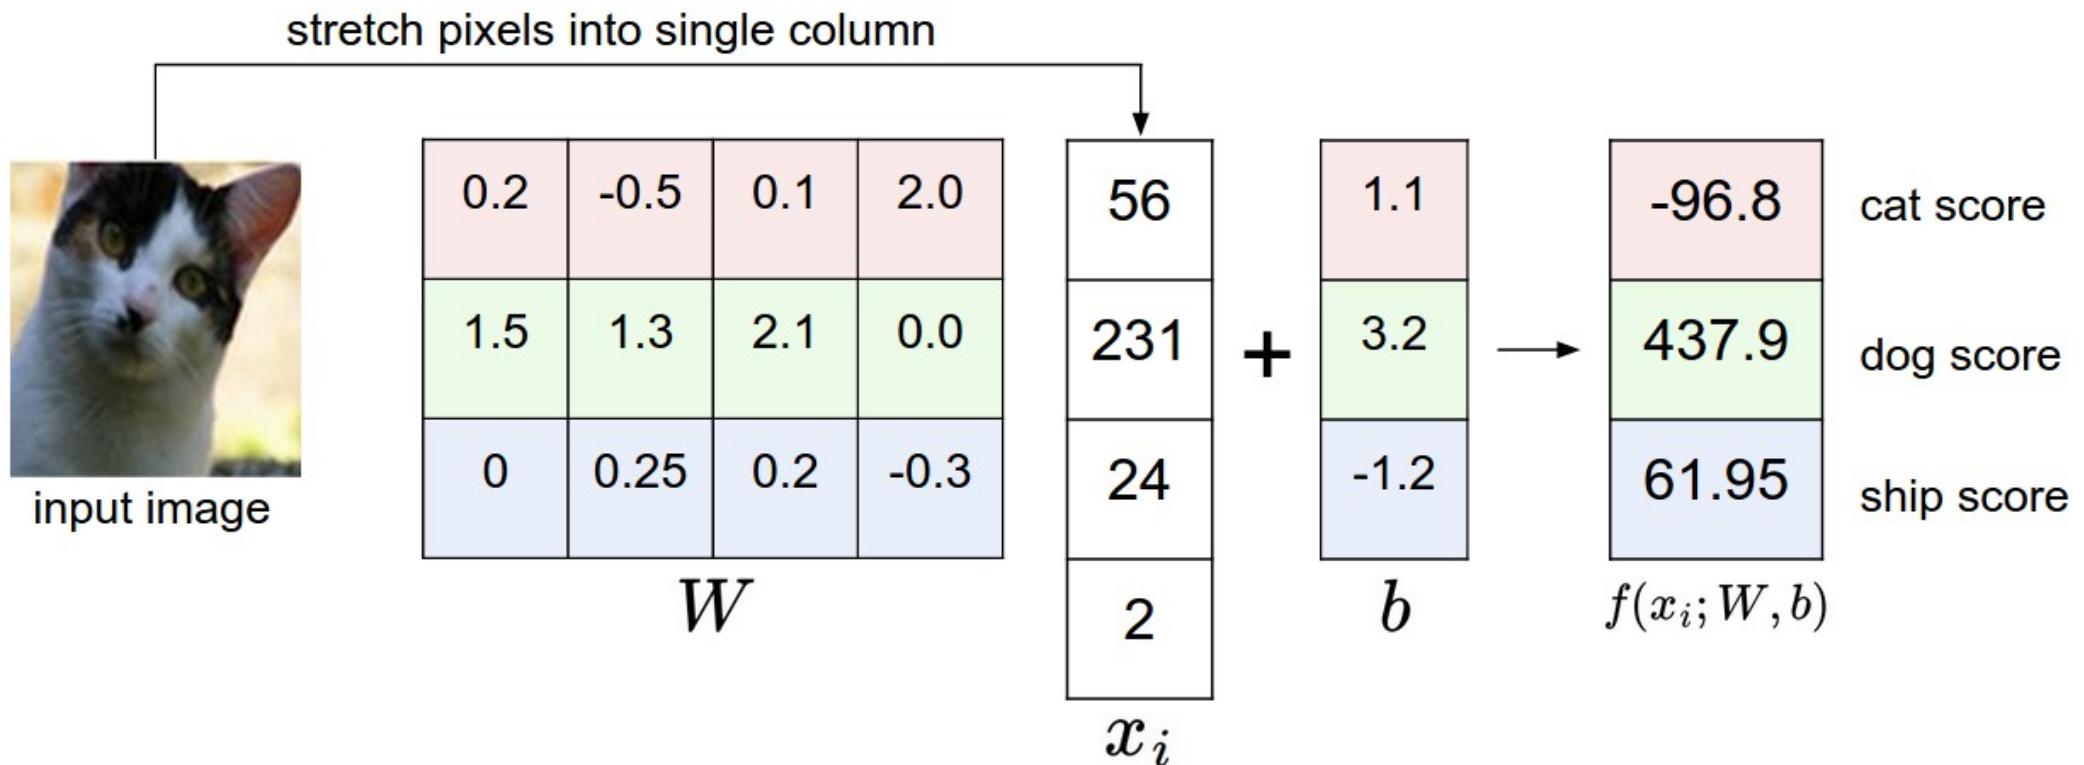
weights  
[10 x 3072]

bias vector  
[10 x 1]



# Linear Classification

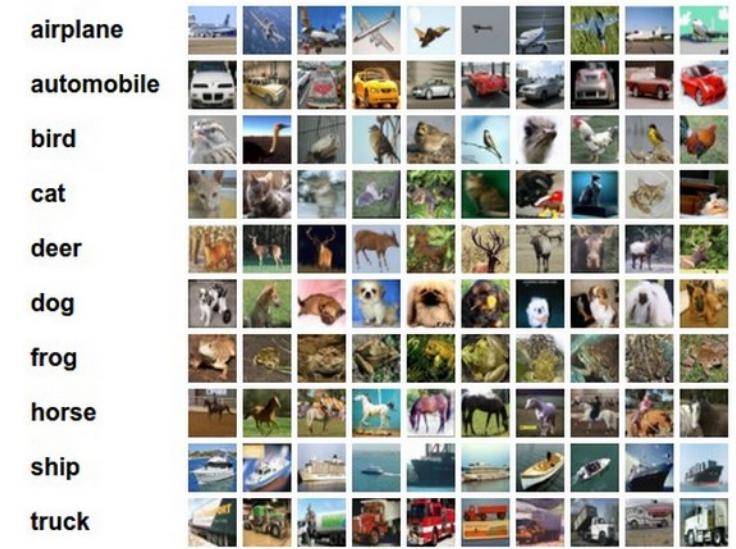
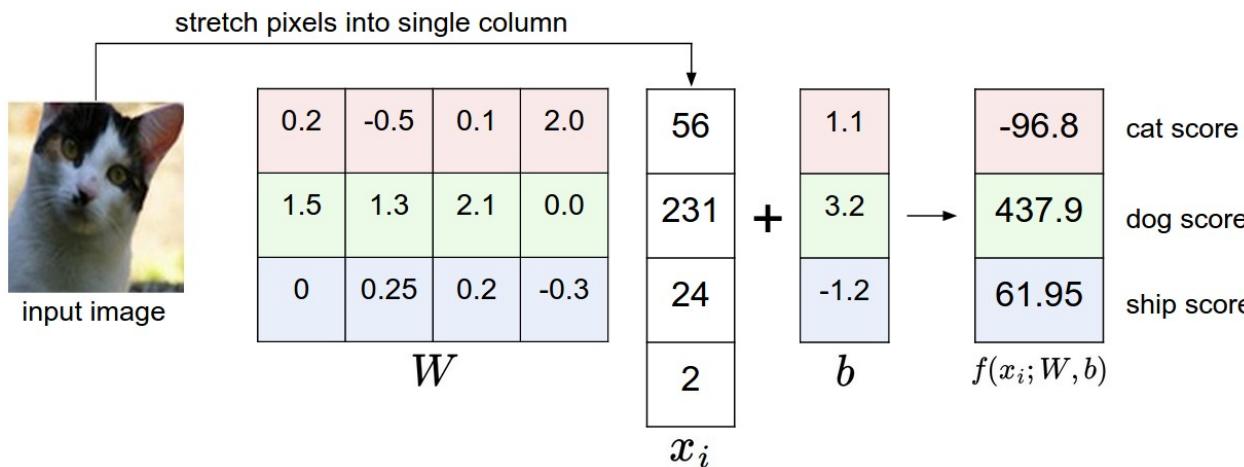
$$f(x_i, W, b) = Wx_i + b$$



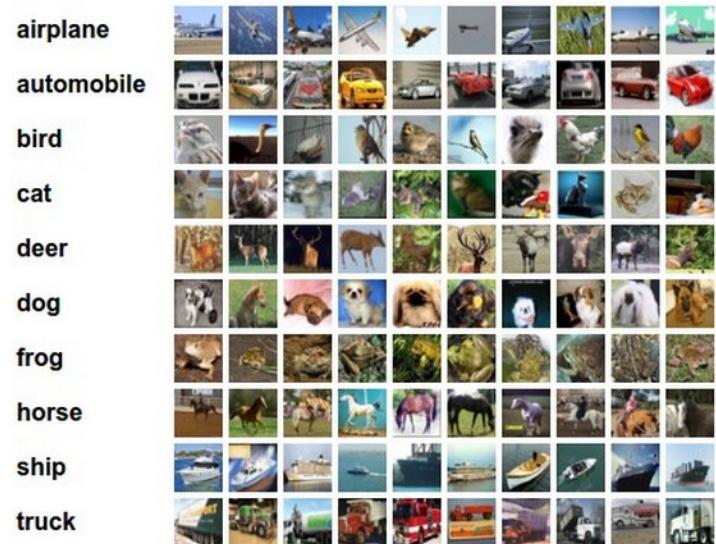
# Interpreting a Linear Classifier

Question:  
what can a linear classifier do?

$$f(x_i, W, b) = Wx_i + b$$



# Interpreting a Linear Classifier

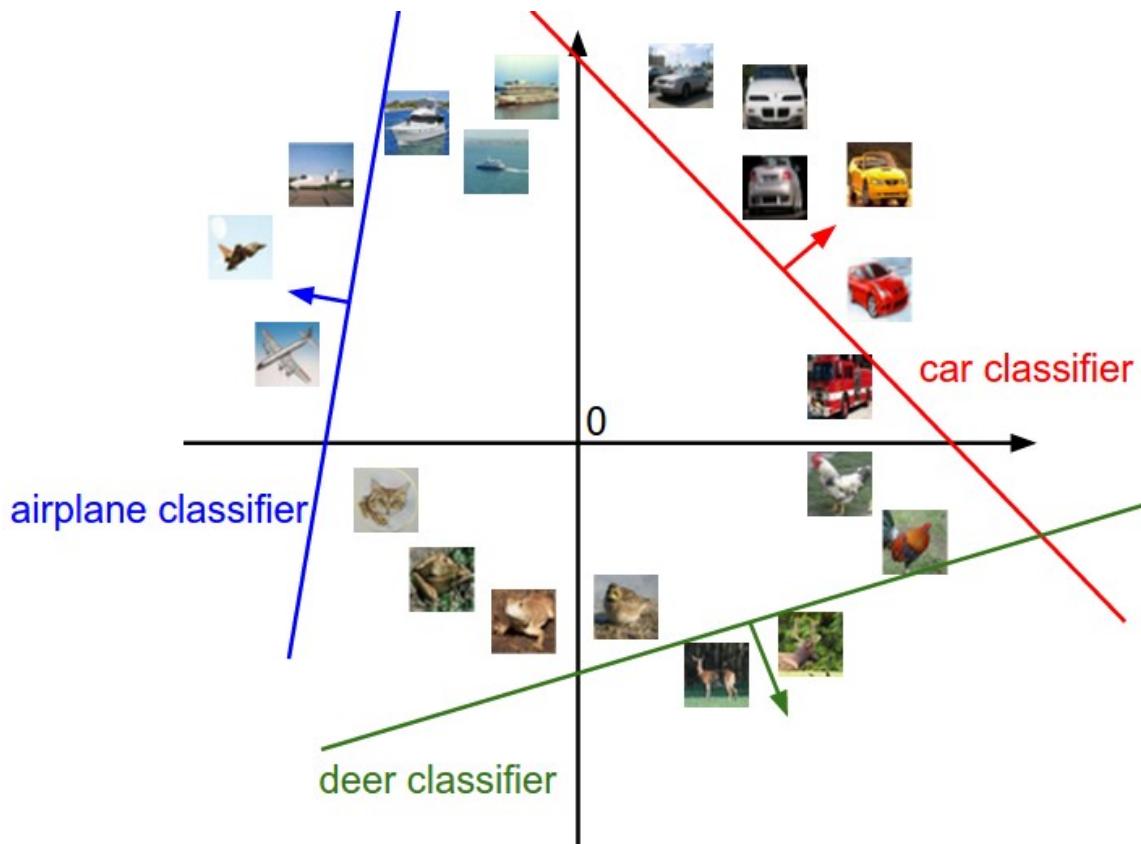


$$f(x_i, W, b) = Wx_i + b$$

Example training  
classifiers on CIFAR-10:



# Interpreting a Linear Classifier



$$f(x_i, W, b) = Wx_i + b$$

# Bias trick

$$f(x_i, W, b) = Wx_i + b \longrightarrow f(x_i, W) = Wx_i$$

0.2	-0.5	0.1	2.0
1.5	1.3	2.1	0.0
0	0.25	0.2	-0.3

$W$

$+ \quad b$

56	1.1
231	3.2
24	-1.2
2	

$x_i$

$\longleftrightarrow$

0.2	-0.5	0.1	2.0	1.1
1.5	1.3	2.1	0.0	3.2
0	0.25	0.2	-0.3	-1.2

$W$

$b$

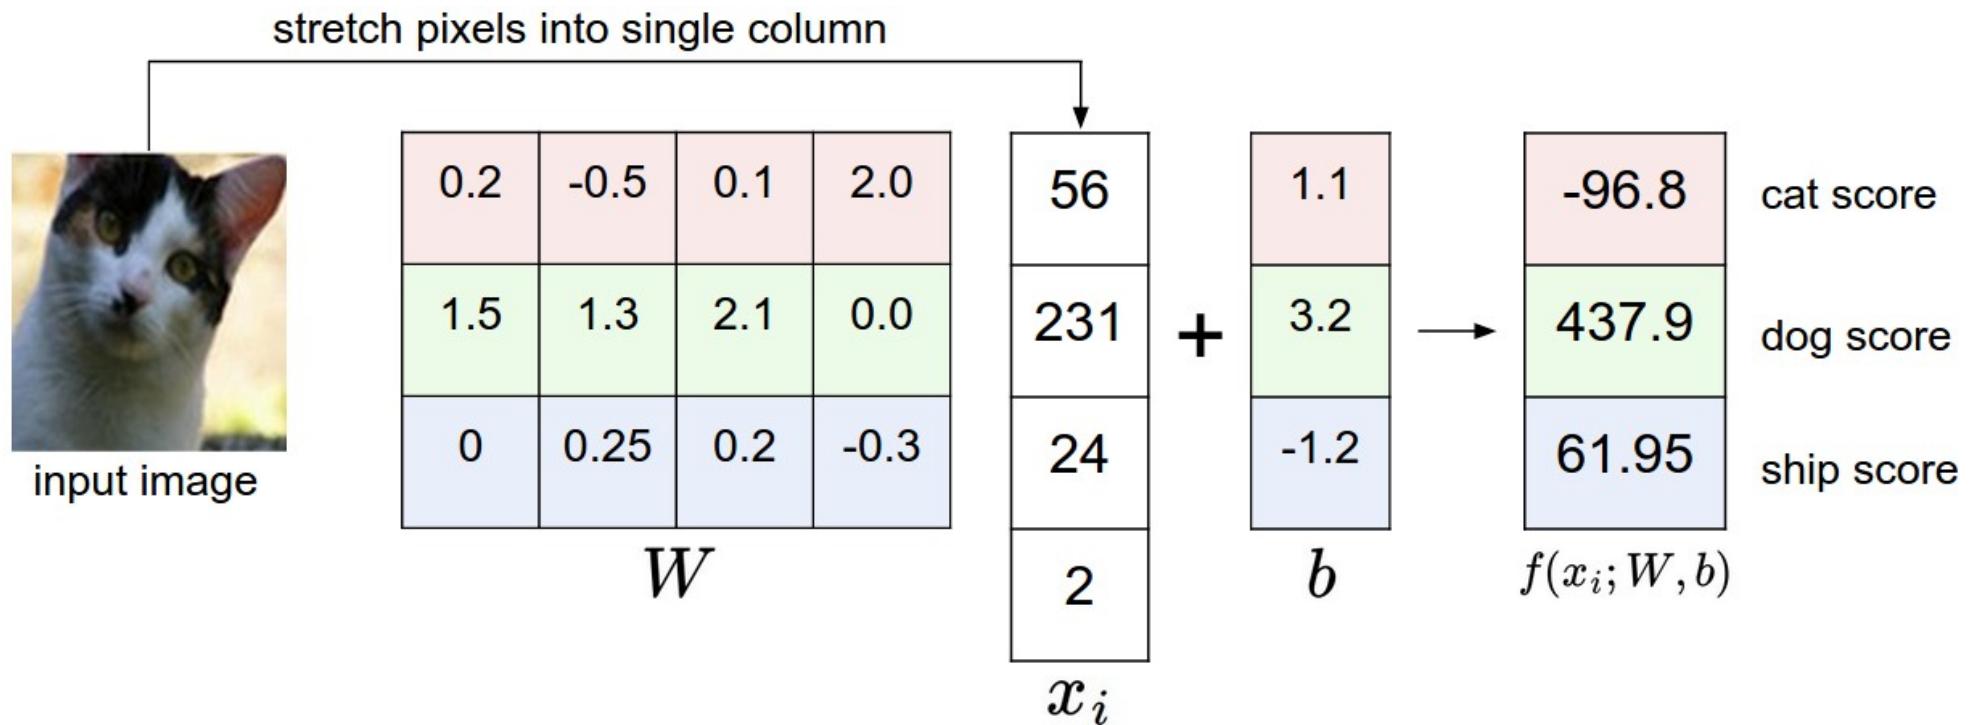
new, single  $W$

56
231
24
2
1

$x_i$

# So far:

We defined a (linear) score function:  $f(x_i, W, b) = Wx_i + b$



## 2. Define a loss function (or cost function, or objective)

## 2. Define a loss function (or cost function, or objective)

- scores, label  $\longrightarrow$  loss.

$$f(x_i, W) \quad y_i \quad L_i$$

## 2. Define a loss function (or cost function, or objective)

- scores, label  $\longrightarrow$  loss.

$$f(x_i, W) \quad y_i \quad L_i$$

---

**Example:**

$$f(x_i, W) = [13, -7, 11]$$

$$y_i = 0$$

## 2. Define a loss function (or cost function, or objective)

- scores, label  $\longrightarrow$  loss.  
 $f(x_i, W)$        $y_i$        $L_i$
- 

**Example:**

$$f(x_i, W) = [13, -7, 11]$$

$$y_i = 0$$

**Question:** if you were to assign a single number to how “unhappy” you are with these scores, what would you do?

2. Define a loss function (or cost function, or objective)  
One (of many ways) to do it: **Multiclass SVM Loss**

$$L_i = \sum_{j \neq y_i} \max(0, f(x_i, W)_j - f(x_i, W)_{y_i} + \Delta)$$

2. Define a loss function (or cost function, or objective)  
One (of many ways) to do it: **Multiclass SVM Loss**

$$L_i = \sum_{j \neq y_i} \max(0, f(x_i, W)_j - f(x_i, W)_{y_i} + \Delta)$$

(One possible generalization of Binary Support Vector Machine to multiple classes)

$$L_i = C \max(0, 1 - y_i w^T x_i) + R(W)$$

## 2. Define a loss function (or cost function, or objective) One (of many ways) to do it: **Multiclass SVM Loss**

$$L_i = \sum_{j \neq y_i} \max(0, f(x_i, W)_j - f(x_i, W)_{y_i} + \Delta)$$

↑  
loss due to  
example i

↑  
sum over all  
incorrect labels

↑  
difference between the correct class  
score and incorrect class score

$$L_i = \sum_{j \neq y_i} \max(0, f(x_i, W)_j - f(x_i, W)_{y_i} + \Delta)$$

loss due to example i      sum over all incorrect labels      difference between the correct class score and incorrect class score



Example:  $L_i = \sum_{j \neq y_i} \max(0, f(x_i, W)_j - f(x_i, W)_{y_i} + \Delta)$

$f(x_i, W) = [13, -7, 11]$  e.g. 10

$y_i = 0$

---

loss = ?

Example:  $L_i = \sum_{j \neq y_i} \max(0, f(x_i, W)_j - f(x_i, W)_{y_i} + \Delta)$

$f(x_i, W) = [13, -7, 11]$  e.g. 10

$y_i = 0$

---

$$L_i = \max(0, -7 - 13 + 10) + \max(0, 11 - 13 + 10)$$

$$L_i = \sum_{j \neq y_i} \max(0, f(x_i, W)_j - f(x_i, W)_{y_i} + \Delta)$$

```
def L_i_vectorized(x, y, W):
    scores = W.dot(x)
    margins = np.maximum(0, scores - scores[y] + 1)
    margins[y] = 0
    loss_i = np.sum(margins)
    return loss_i
```

There is a bug with the objective...



$$L = \frac{1}{N} \sum_i \sum_{j \neq y_i} \left[ \max(0, f(x_i; W)_j - f(x_i; W)_{y_i} + \Delta) \right]$$

# L2 Regularization

$$R(W) = \sum_k \sum_l W_{k,l}^2$$

$$L = \frac{1}{N} \sum_i \sum_{j \neq y_i} \left[ \max(0, f(x_i; W)_j - f(x_i; W)_{y_i} + \Delta) \right] + \boxed{\lambda R(W)}$$

Regularization strength

## L2 regularization: motivation

$$\mathbf{x} = [1, 1, 1, 1]$$

$$w_1 = [1, 0, 0, 0]$$

$$w_2 = [0.25, 0.25, 0.25, 0.25]$$

$$w_1^T \mathbf{x} = w_2^T \mathbf{x} = 1$$

$$L = \frac{1}{N} \sum_i \sum_{j \neq y_i} \left[ \max(0, f(x_i; W)_j - f(x_i; W)_{y_i} + \Delta) \right] + \lambda R(W)$$



$$L = \frac{1}{N} \sum_i \sum_{j \neq y_i} \left[ \max(0, f(x_i; W)_j - f(x_i; W)_{y_i} + \Delta) \right] + \lambda R(W)$$

Do we have to cross-validate both  
 $\Delta$  and  $\lambda$  ?

$$L = \frac{1}{N} \sum_i \sum_{j \neq y_i} \left[ \max(0, f(x_i; W)_j - f(x_i; W)_{y_i}) + \boxed{\Delta} \right] + \lambda R(W)$$

$$\begin{aligned} & \max(0, \delta f + \Delta) \\ \rightarrow & \max(0, \delta f + \Delta_2) && \text{if } \Delta \rightarrow \Delta_2 \\ \rightarrow & \max(0, \frac{\Delta_2}{\Delta} \delta f + \Delta_2) && \text{then scale the weights by } \frac{\Delta_2}{\Delta} \\ = & \max(0, \frac{\Delta_2}{\Delta} (\delta f + \Delta)) \\ = & \frac{\Delta_2}{\Delta} \max(0, \delta f + \Delta) \end{aligned}$$

So far...

## 1. Score function

$$f(x_i, W, b) = Wx_i + b$$

## 2. Loss function

$$L = \frac{1}{N} \sum_i \sum_{j \neq y_i} \left[ \max(0, f(x_i; W)_j - f(x_i; W)_{y_i} + \Delta) \right] + \lambda R(W)$$

