


Brewing Deep Neural Networks with Caffe



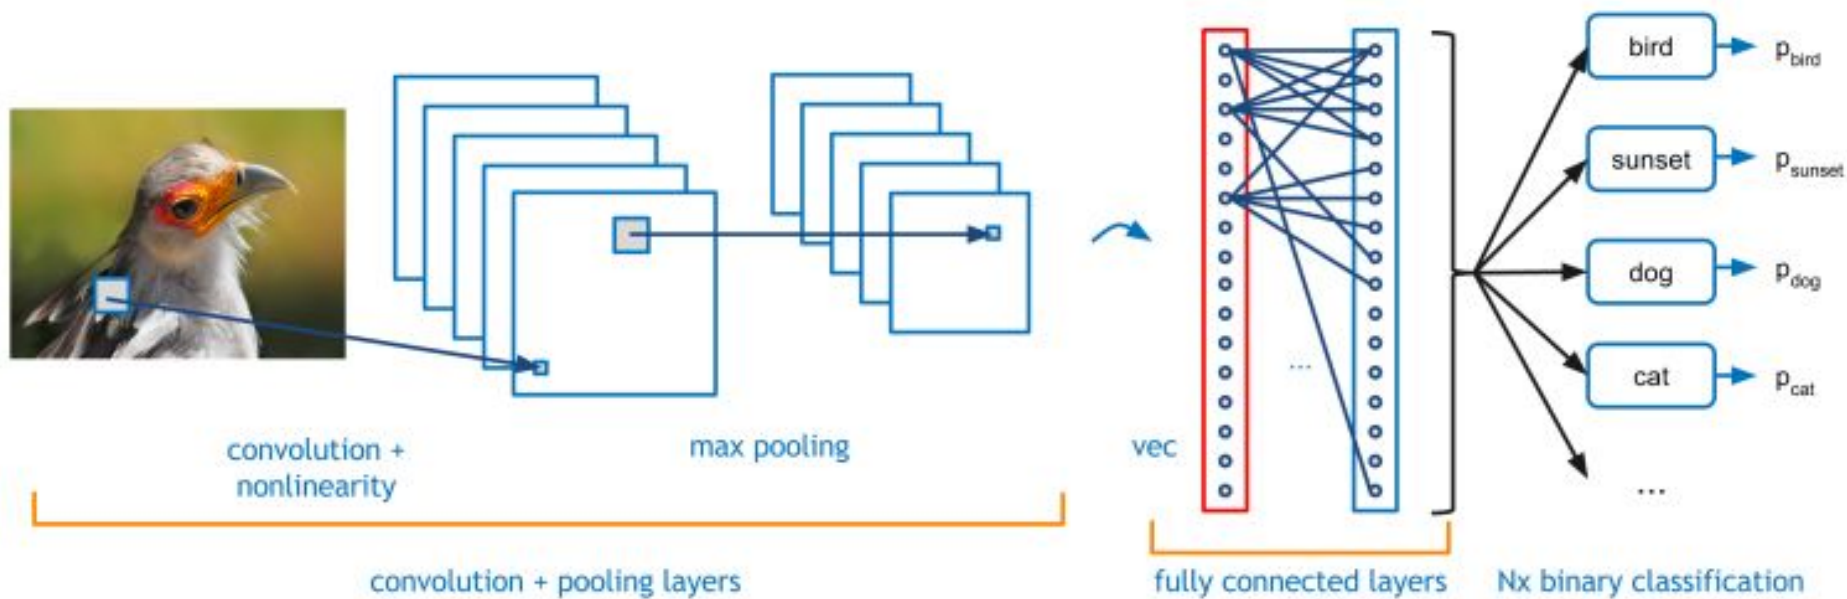
Maximally accurate	Maximally specific
espresso	2.23192
coffee	2.19914
beverage	1.93214
liquid	1.89367
fluid	1.85519

Abhimanyu Dubey

ELL881/ELL784

Slides at <http://iitd.info/hpccaffe>


Adapted from original caffe slides




All in a day's work with Caffe

Reference Models


AlexNet: ImageNet Classification



R-CNN: Regions with CNN features



GoogLeNet: ILSVRC14 winner



Caffe offers the

- model definitions
- optimization settings
- pre-trained weights

so you can start right away.

The BVLC models are licensed for unrestricted use.

The community shares models in the [Model Zoo](#).

Net

- A network is a set of layers and their connections:

name: "dummy-net"

layer { name: "data" ...}


layer { name: "conv" ...}

layer { name: "pool" ...}

... more layers ...

layer { name: "loss" ...}


- Caffe creates and checks the net from the definition.
- Data and derivatives flow through the net as *blobs* – an array interface



LogReg ↑


LeNet →

ImageNet, Krizhevsky 2012 →




DAG

Many current deep models have linear structure




but Caffe nets can have any directed acyclic graph (DAG) structure.


Define bottoms and tops
and Caffe will connect the net.



SDS two-stream net



GoogLeNet Inception Module



LRCN joint vision-sequence model

Layer Protocol

Setup: run once for initialization.

Forward: make output given input.


Backward: make gradient of output

- w.r.t. bottom
- w.r.t. parameters (if needed)

Reshape: set dimensions.

Compositional Modeling

The Net's forward and backward passes are composed of the layers' steps.



[Layer Development Checklist](#)

Protobuf Model Format

- Strongly typed format
- Auto-generates code
- Developed by Google
- Defines Net / Layer / Solver

schemas in **caffe.proto**

```
message ConvolutionParameter {  
    // The number of outputs for the layer  
    optional uint32 num_output = 1;  
    // whether to have bias terms  
    optional bool bias_term = 2 [default = true];  
}
```

```
name: "conv1"  
type: "Convolution"  
bottom: "data"  
top: "conv1"  
convolution_param {  
    num_output: 20  
    kernel_size: 5  
    stride: 1  
    weight_filler {  
        type: "xavier"  
    }  
}
```

Model Zoo Format

readme.md

Raw

name: FCN-32s Fully Convolutional Semantic Segmentation on PASCAL-Context caffemodel: fcn-32s-pascalcontext.caffemodel caffemodel_url: <http://dl.caffe.berkeleyvision.org/fcn-32s-pascalcontext.caffemodel> sha1: adbbd504c280e2b8966fc32e32ada2a2ecf13603

gist_id: 80667189b218ad570e82

This is a model from the [paper](#):

```
Fully Convolutional Networks for Semantic Segmentation
Jonathan Long, Evan Shelhamer, Trevor Darrell
arXiv:1411.4038
```

Gists on github hold model definition, license, url for weights, and hash of Caffe commit that guarantees compatibility.

Solving: Training a Net

Optimization like model definition is configuration.

`train_net: "lenet_train.prototxt"`

`base_lr: 0.01`

`momentum: 0.9`

`weight_decay: 0.0005`

`max_iter: 10000`

`snapshot_prefix: "lenet_snapshot"`

All you need to run things
on the GPU.



`> caffe train -solver lenet_solver.prototxt -gpu 0`

Stochastic Gradient Descent (SGD) + momentum ·

Adaptive Gradient (ADAGRAD) · Nesterov's Accelerated Gradient (NAG)

Solver Showdown: MNIST Autoencoder

AdaGrad

```
I0901 13:36:30.007884 24952 solver.cpp:232] Iteration 65000, loss = 64.1627
I0901 13:36:30.007922 24952 solver.cpp:251] Iteration 65000, Testing net (#0) # train set
I0901 13:36:33.019305 24952 solver.cpp:289] Test loss: 63.217
I0901 13:36:33.019356 24952 solver.cpp:302]      Test net output #0: cross_entropy_loss = 63.217 (* 1 = 63.217 loss)
I0901 13:36:33.019773 24952 solver.cpp:302]      Test net output #1: l2_error = 2.40951
```

SGD

```
I0901 13:35:20.426187 20072 solver.cpp:232] Iteration 65000, loss = 61.5498
I0901 13:35:20.426218 20072 solver.cpp:251] Iteration 65000, Testing net (#0) # train set
I0901 13:35:22.780092 20072 solver.cpp:289] Test loss: 60.8301
I0901 13:35:22.780138 20072 solver.cpp:302]      Test net output #0: cross_entropy_loss = 60.8301 (* 1 = 60.8301 loss)
I0901 13:35:22.780146 20072 solver.cpp:302]      Test net output #1: l2_error = 2.02321
```

Nesterov

```
I0901 13:36:52.466069 22488 solver.cpp:232] Iteration 65000, loss = 59.9389
I0901 13:36:52.466099 22488 solver.cpp:251] Iteration 65000, Testing net (#0) # train set
I0901 13:36:55.068370 22488 solver.cpp:289] Test loss: 59.3663
I0901 13:36:55.068410 22488 solver.cpp:302]      Test net output #0: cross_entropy_loss = 59.3663 (* 1 = 59.3663 loss)
I0901 13:36:55.068418 22488 solver.cpp:302]      Test net output #1: l2_error = 1.79998
```

Weight Sharing

- Just give the parameter blobs explicit names using the `param` field
- Layers specifying the same `param` name will share that parameter, accumulating gradients accordingly

```
layer: {  
  name: 'innerproduct1'  
  type: INNER_PRODUCT  
  inner_product_param {  
    num_output: 10  
    bias_term: false  
    weight_filler {  
      type: 'gaussian'  
      std: 10  
    }  
  }  
  param: 'sharedweights'  
  bottom: 'data'  
  top: 'innerproduct1'  
}  
layer: {  
  name: 'innerproduct2'  
  type: INNER_PRODUCT  
  inner_product_param {  
    num_output: 10  
    bias_term: false  
  }  
  param: 'sharedweights'  
  bottom: 'data'  
  top: 'innerproduct2'  
}
```

Recipe for Brewing

- Convert the data to Caffe-format
 - lmdb, leveldb, hdf5 / .mat, list of images, etc.
- Define the Net
- Configure the Solver
- `caffe train -solver solver.prototxt -gpu 0`
- Examples are your friends
 - `caffe/examples/mnist, cifar10, imagenet`
 - `caffe/examples/*.ipynb`
 - `caffe/models/*`

Brewing Models

from logistic regression to non-linearity

[see notebook](#)

Take a pre-trained model and fine-tune to new tasks


[DeCAF] [Zeiler-Fergus] [OverFeat]

Lots of Data



image by Andrej Karpathy

Your Task



**Style
Recognition**



© kaggle.com

**Dogs vs.
Cats**
top 10 in
10 minutes

From ImageNet to Style

Simply change a few lines in the model definition

```
layer {
  name: "data"
  type: "Data"
  data_param {
    source: "ilsvrc12_train_lmdb"
    mean_file: "../../data/ilsvrc12"
    ...
  }
  ...
}
...
layer {
  name: "fc8"
  type: "InnerProduct"
  inner_product_param {
    num_output: 1000
    ...
  }
}
```

```
layer {
  name: "data"
  type: "Data"
  data_param {
    source: "style_train_lmdb"
    mean_file: "../../data/ilsvrc12"
    ...
  }
  ...
}
...
layer {
  name: "fc8-style"
  type: "InnerProduct"
  inner_product_param {
    num_output: 20
    ...
  }
}
```

Input:
A different source

new name =
new params

Last Layer:
A different classifier

From ImageNet to Style

```
> caffe train -solver models/finetune_flickr_style/solver.prototxt  
              -weights bvlc_reference_caffenet.caffemodel
```

Step-by-step in pycaffe:

```
pretrained_net = caffe.Net(  
    "net.prototxt", "net.caffemodel")  
solver = caffe.SGDSolver("solver.prototxt")  
solver.net.copy_from(pretrained_net)  
solver.solve()
```

Vintage HDR Melancholy Minimal



Fine-tuning

transferring features to style recognition

[see notebook](#)


When to Fine-tune?

A good first step

- More robust optimization – good initialization helps
- Needs less data
- Faster learning

State-of-the-art results in

- recognition
- detection
- segmentation



[Zeiler-Fergus]

Fine-tuning Tricks

Learn the last layer first

- Caffe layers have local learning rates: `param { lr_mult: 1 }`
- Freeze all but the last layer for fast optimization and avoiding early divergence by setting `lr_mult: 0` to fix a parameter.
- Stop if good enough, or keep fine-tuning


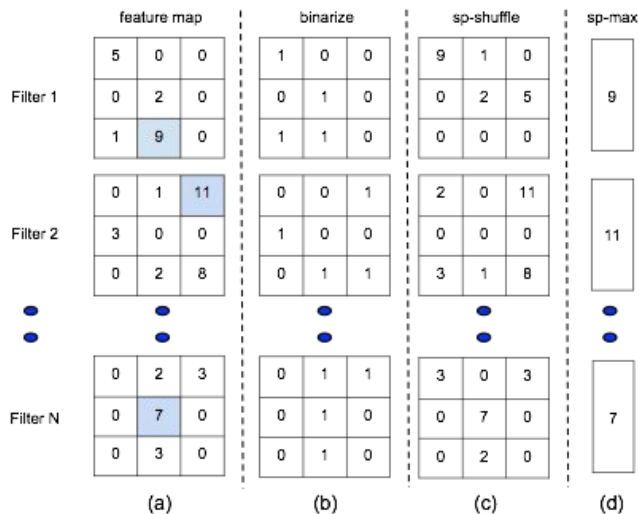
Reduce the learning rate

- Drop the solver learning rate by 10x, 100x
- Preserve the initialization from pre-training and avoid divergence

Do net surgery see notebook on [editing model parameters](#)

After fine-tuning

- Supervised pre-training does not overfit
- Representation is (mostly) distributed
- Sparsity comes “for free” in deep representation



Running Caffe on HPC

HPC

- High performance distributed computing system (<http://supercomputing.iitd.ac.in>)
- 422 compute nodes, 161 of which are GPU nodes
- 4 general (CPU) login nodes and 2 GPU login nodes

HPC continued

- Login nodes aren't meant for running jobs - **so please don't** (they'll be killed anyway)
- Jobs are handled by a scheduling system known as PBS (Portable Batch System)
- PBS handles the distribution and scheduling of jobs (more at <http://supercomputing.iitd.ac.in/?pbs>)

HPC continued

- Two storage quotas
 - 30 GB on home (~)
 - 200 TB on scratch (/scratch)
- Hack to make storage simpler
 - Create a symlink from your scratch folder to a folder in your home directory
 - `ln -s /scratch/ee/btech/ee2110061/
/home/ee/btech/ee2110061/scratch`

HPC continued

- Keep all datasets, models and logs in scratch (duh)
- Use screen (<https://www.rackaid.com/blog/linux-screen-tutorial-and-how-to/>) to keep your sessions alive
- Don't use nohup - very easy to forget to kill and then eventually use up all disk space with a huge nohup file
- To have internet access, proxy login is required -
 - lynx is a CLI browser you can use

SSH keys

- Required only for passwordless login - that's it
- Straightforward steps:
 - On your laptop/host machine:
 - `ssh-keygen -t rsa` #(keep pressing enter)
 - `ssh-add`
 - `cat ~/.ssh/id_rsa.pub`
 - #Copy EVERYTHING that is outputted
 - On HPC (after login - yes, using your password)
 - `echo "<paste-here>" >> ~/.ssh/authorized_keys`
 - Logout and login again without a password

Caffe on HPC

- Caffe's available as a module
(<http://modules.sourceforge.net>)
- Caffe source isn't accessible on HPC - only executables and tests are kept - but that's all you need ATM
- Remember to load Caffe as part of the job submission script

Caffe on HPC

- To load Caffe - module load apps/Caffe (small c for CPU-only - apps/caffe)
- Check where Caffe executables are by running
 - echo \$CAFFE_ROOT/bin
- You can try out Caffe by running the examples - you will have to modify the training scripts though

pyCaffe

- Pycaffe is loaded at \$CAFFE_ROOT/python
- To use “import caffe” on HPC, run
 - mkdir -p /home/ee/btech/ee2110061/.local/lib/python2.7/site-packages/ #(if you’ve not done this before)
 - ln -s \$CAFFE_ROOT/python /home/ee/btech/ee2110061/.local/lib/python2.7/site-packages/caffe